

---

# **DDBMS Security**

# Introduction

---

- Data security

- Protect data against unauthorized access.

- Two aspects

- Data protection.

- Authorization Control.

# Aspects of Data security

---

## ■ Data Protection

- Can be achieved using data encryption techniques.

## ■ Authorization Control

- It ensures that only authorized users perform, operations that they are allowed to perform on the database.

# Authorization Control

---

- It includes two main issues

- Access control

- ◆ Unauthorized Access to data should not be allowed.

- Integrity

- ◆ Only authorized users should be allowed to modify data in the database.

# Semantic Data Control

---

## ■ Involves:

- » View management
- » Security control
- » Integrity control

## ■ Objective :

- » Insure that **authorized** users perform **correct** operations on the database, contributing to the maintenance of the database integrity.

# View Management

---

View – virtual relation

- generated from base relation(s) by a query
- not stored as base relations
- Stored as a definition

Example :

```
CREATE VIEW    SYSAN (ENO, ENAME)
AS           SELECT ENO, ENAME
               FROM    EMP
               WHERE  TITLE="Syst. Anal."
```

EMP

ENO	ENAME	TITLE
E1	J. Doe	Elect. Eng
E2	M. Smith	Syst. Anal.
E3	A. Lee	Mech. Eng.
E4	J. Miller	Programmer
E5	B. Casey	Syst. Anal.
E6	L. Chu	Elect. Eng.
E7	R. Davis	Mech. Eng.
E8	J. Jones	Syst. Anal.

SYSAN

ENO	ENAME
E2	M.Smith
E5	B.Casey
E8	J.Jones

# View Management

---

Views can be manipulated as base relations

Example :

```
SELECT   ENAME, PNO, RESP
FROM     SYSAN, ASG
WHERE     SYSAN.ENO = ASG.ENO
```

# Query Modification

---

queries expressed on views



queries expressed on base relations

Example :

```
SELECT ENAME, PNO, RESP
FROM    SYSAN, ASG
WHERE   SYSN.ENO = ASG.ENO
```



```
SELECT ENAME, PNO, RESP
FROM    EMP, ASG
WHERE   EMP.ENO = ASG.ENO
AND     TITLE = "Syst. Anal."
```

ENAME	PNO	RESP
M.Smith	P1	Analyst
M.Smith	P2	Analyst
B.Casey	P3	Manager
J.Jones	P4	Manager



# View Management

---

- To restrict access

```
CREATE    VIEW      ESAME
AS       SELECT    *
          FROM      EMP E1, EMP E2
          WHERE     E1.TITLE = E2.TITLE
          AND       E1.ENO = USER
```

# View Management in DDBMS

---

- Views might be derived from fragments.
- View definition storage should be treated as database storage
- View might be costly if base relations are distributed
  - » use **snapshots (Materialized View)**
    - ◆ Static views - do not reflect the updates to the base relations
    - ◆ managed as temporary relations - only access path is sequential scan
    - ◆ bad selectivity - snapshots behave as pre-calculated answers
    - ◆ periodic recalculation

# Data Security

---

## ■ Data protection

- ▶▶▶ Prevents the physical content of data to be *understood* by unauthorized users
- ▶▶▶ Uses encryption/decryption techniques (Public key)

## ■ Access control

- ▶▶▶ Only authorized users perform operations they are allowed to on database objects
- ▶▶▶ Discretionary access control (DAC)
  - ◆ Long been provided by DBMS with authorization rules
- ▶▶▶ Multilevel access control (MAC)
  - ◆ Increases security with security levels

# Methods of Access Control

---

- Discretionary Access Control

- Based on privileges or access rights

- Mandatory Access control

- Based on policies that can't be changed by individual users

# Discretionary Access Control

---

## ■ Main actors

- ▶▶▶ Subjects (users, groups of users) who execute operations
- ▶▶▶ Operations (in queries or application programs)
- ▶▶▶ Objects, on which operations are performed

## ■ Checking whether a subject may perform an operation on an object

- ▶▶▶ Authorization= (subject, op. type, object def.)
- ▶▶▶ Defined using **GRANT OR REVOKE**
- ▶▶▶ **Centralized**: one single user class (admin.) may grant or revoke
- ▶▶▶ **Decentralized**, with op. type GRANT
  - ◆ More flexible but recursive revoking process which needs the hierarchy of grants

# Multilevel Access Control

---

- Different security levels (*clearances*)

- *Top Secret > Secret > Confidential > Unclassified*

- Access controlled by 2 rules:

- No read up

- ◆ subject  $S$  is allowed to read an object of level  $L$  only if  $level(S) \geq L$
    - ◆ Protect data from unauthorized disclosure, e.g. **a subject with secret clearance cannot read top secret data**

- No write down:

- ◆ subject  $S$  is allowed to write an object of level  $L$  only if  $level(S) \leq L$
    - ◆ Protect data from unauthorized change, e.g. **a subject with top secret clearance can only write top secret data but not secret data** (which could then contain top secret data)

# Semantic Integrity Control

---

- Maintain database consistency by enforcing a set of constraints defined on the database.
- Structural constraints
  - basic semantic properties inherent to a data model  
e.g., unique key constraint in relational model
- Behavioral constraints
  - regulate application behavior
    - ◆ e.g., dependencies in the relational model
- Two components
  - Integrity constraint specification
  - Integrity constraint enforcement

# Semantic Integrity Control

---

## ■ Procedural

- control embedded in each application program

## ■ Declarative

- assertions in predicate calculus
- easy to define constraints
- definition of database consistency clear
- inefficient to check assertions for each update
  - ◆ limit the search space
  - ◆ decrease the number of data accesses/assertion
  - ◆ preventive strategies
  - ◆ checking at compile time



# Constraint Specification Language

---

## Predefined constraints

specify the more common constraints of the relational model

### ►►► Not-null attribute

ENO **NOT NULL IN** EMP

### ►►► Unique key

(ENO, PNO) **UNIQUE IN** ASG

### ►►► Foreign key

A key in a relation  $R$  is a foreign key if it is a primary key of another relation  $S$  and the existence of any of its values in  $R$  is dependent upon the existence of the same value in  $S$

PNO **IN** ASG **REFERENCES** PNO **IN** PROJ

### ►►► Functional dependency

ENO **IN** EMP **DETERMINES** ENAME

# Constraint Specification Language

---

## Precompiled constraints

Express preconditions that must be satisfied by all tuples in a relation for a given update type

(INSERT, DELETE, MODIFY)

NEW - ranges over new tuples to be inserted

OLD - ranges over old tuples to be deleted

General Form

```
CHECK ON <relation> [WHEN <update type>]  
    <qualification>
```

# Constraint Specification Language

---

## Precompiled constraints

### Domain constraint

```
CHECK ON PROJ (BUDGET ≥ 500000 AND BUDGET ≤ 1000000)
```

### Domain constraint on deletion

```
CHECK ON PROJ WHEN DELETE (BUDGET = 0)
```

### Transition constraint

```
CHECK ON PROJ (NEW.BUDGET > OLD.BUDGET AND  
                  NEW.PNO = OLD.PNO)
```

# Constraint Specification Language

---

## General constraints

Constraints that must always be true. Formulae of tuple relational calculus where all variables are quantified.

### General Form

**CHECK ON** <variable>:<relation>, (<qualification>)

#### Functional dependency

**CHECK ON** e1:EMP, e2:EMP

(e1.ENAME = e2.ENAME **IF** e1.ENO = e2.ENO)

#### Constraint with aggregate function

**CHECK ON** g:ASG, j:PROJ

(**SUM**(g.DUR **WHERE** g.PNO = j.PNO) < 100 **IF**  
j.PNAME = "CAD/CAM")

# Integrity Enforcement

---

Two methods

## ■ Detection

Execute update  $u: D \rightarrow D_u$

If  $D_u$  is inconsistent then

    compensate  $D_u \rightarrow D_u'$

else

    undo  $D_u \rightarrow D$

## ■ Preventive

Execute  $u: D \rightarrow D_u$  only if  $D_u$  will be consistent

    ▶▶▶ Determine valid programs

    ▶▶▶ Determine valid states

# Query Modification

---

- Preventive
- Add the assertion qualification to the update query
- Only applicable to tuple calculus formulae with universally quantified variables

```
UPDATE PROJ
SET      BUDGET = BUDGET*1.1
WHERE    PNAME ="CAD/CAM"
      ↓
UPDATE PROJ
SET      BUDGET = BUDGET*1.1
WHERE    PNAME ="CAD/CAM"
AND      NEW.BUDGET ≥ 500000
AND      NEW.BUDGET ≤ 1000000
```

# Compiled Assertions

---

Triple  $(R, T, C)$  where

$R$  relation

$T$  update type (insert, delete, modify)

$C$  assertion on differential relations

Example: Foreign key assertion

$\forall g \in \text{ASG}, \exists j \in \text{PROJ} : g.\text{PNO} = j.\text{PNO}$

Compiled assertions:

$(\text{ASG}, \text{INSERT}, C1), (\text{PROJ}, \text{DELETE}, C2), (\text{PROJ}, \text{MODIFY}, C3)$

where

$C1: \forall \text{NEW} \in \text{ASG}^+, \exists j \in \text{PROJ} : \text{NEW.PNO} = j.\text{PNO}$

$C2: \forall g \in \text{ASG}, \forall \text{OLD} \in \text{PROJ}^- : g.\text{PNO} \neq \text{OLD.PNO}$

$C3: \forall g \in \text{ASG}, \forall \text{OLD} \in \text{PROJ}^-, \exists \text{NEW} \in \text{PROJ}^+ : g.\text{PNO} \neq \text{OLD.PNO}$   
**OR OLD.PNO = NEW.PNO**

# Differential Relations

---

Given relation  $R$  and update  $u$

$R^+$  contains tuples inserted by  $u$

$R^-$  contains tuples deleted by  $u$

Type of  $u$

insert  $R^-$  empty

delete  $R^+$  empty

modify  $R^+ \cup (R - R^-)$



# Differential Relations

---

## Algorithm

- Input:** Relation  $R$ , update  $u$ , compiled assertion  $C_i$
- Step 1:** Generate differential relations  $R^+$  and  $R^-$
- Step 2:** Retrieve the tuples of  $R^+$  and  $R^-$  which **do not** satisfy  $C_i$
- Step 3:** If retrieval is not successful, then the assertion is valid.

## Example :

$u$  is delete on J. Enforcing (J, DELETE, C2) :

*retrieve all* tuples of J-

*into* RESULT

*where* not(C2)

If  $\text{RESULT} = \emptyset$ , the assertion is verified.

# Distributed Integrity Control

---

## ■ Problems:

- Definition of constraints
  - ◆ consideration for fragments
- Where to store
  - ◆ replication
  - ◆ non-replicated : fragments
- Enforcement
  - ◆ minimize costs

# Types of Distributed Assertions

---

- Individual assertions
  - single relation, single variable
  - domain constraint
- Set oriented assertions
  - single relation, multi-variable
    - ◆ functional dependency
  - multi-relation, multi-variable
    - ◆ foreign key
- Assertions involving aggregates

# Distributed Integrity Control

---

## ■ Assertion Definition

- similar to the centralized techniques
- transform the assertions to compiled assertions

## ■ Assertion Storage

- Individual assertions
  - ◆ one relation, only fragments
  - ◆ at each fragment site, check for compatibility
  - ◆ if compatible, store; otherwise reject
  - ◆ if all the sites reject, globally reject
- Set-oriented assertions
  - ◆ involves joins (between fragments or relations)
  - ◆ maybe necessary to perform joins to check for compatibility
  - ◆ store if compatible

# Distributed Integrity Control

---

## ■ Assertion Enforcement

### ➤ Where do you enforce each assertion?

- ◆ type of assertion
- ◆ type of update and where update is issued

### ➤ Individual Assertions

- ◆ update = insert
  - ✓ enforce at the site where the update is issued
- ◆ update = qualified
  - ✓ send the assertions to all the sites involved
  - ✓ execute the qualification to obtain R+ and R-
  - ✓ each site enforce its own assertion

### ➤ Set-oriented Assertions

- ◆ single relation
  - ✓ similar to individual assertions with qualified updates
- ◆ multi-relation
  - ✓ move data between sites to perform joins; then send the result to the query master site